

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import mean_squared_error

from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import mutual_info_regression

import os
os.getcwd()
```

```
Out[1]: 'C:\\Users\\csant\\python_analysis_update'
```

```
In [2]: # Load CSV file of curated/merged data
df = pd.read_csv('c:/csc606/curated colocated data/merged_velocity_data.csv', low_me
```

```
In [3]: # Create single column in dataframe called "key" that concatenates leg, site, hole,
# The key column will be used for joining/merging later
df['key']=df['leg'] + "." + df['site'] + "." + df['hole'] + "." + df['core'] + "."

```

```
In [4]: # Create the rho feature, as recommended by Taylor
# rho = absolute difference between GRA density and 0.31 * velocity ^ 0.25
df['rho']=abs(df['GRA_density(g/cm^3)']-(0.31*df['compressional_velocity(m/s)']**0.
```

```
In [5]: df
```

Out[5]:

| | latitude(dd) | longitude(dd) | leg | site | hole | core | section | depth_m | compr |
|-----------------|--------------|---------------|------|-------|------|------|---------|----------|-------|
| 0 | 31.789787 | 139.026217 | 350 | U1437 | D | 24 | 4 | 645.330 | |
| 1 | 31.789787 | 139.026217 | 350 | U1437 | D | 53 | 1 | 912.683 | |
| 2 | 31.789798 | 139.026523 | 350 | U1437 | E | 5 | 3 | 1113.300 | |
| 3 | -38.829700 | 178.476055 | 372A | U1517 | C | 8 | 2 | 41.092 | |
| 4 | -38.859490 | 178.896032 | 375 | U1518 | F | 30 | 7 | 472.850 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 26157261 | 19.489617 | -82.936100 | 165 | 998 | A | 7 | 5 | 62.610 | |
| 26157262 | 16.553717 | -79.867400 | 165 | 1000 | A | 9 | 3 | 73.100 | |
| 26157263 | 55.477183 | -14.650867 | 162 | 981 | A | 7 | 1 | 55.150 | |
| 26157264 | 1.202300 | -83.737000 | 111 | 677 | A | 27 | 5 | 247.550 | |
| 26157265 | 16.130700 | 60.744000 | 117 | 720 | A | 1 | 4 | 4.900 | |

26157266 rows × 21 columns

In [6]:

```
# Load CSV file of Taylor's Labeled data
labeled_df = pd.read_csv('c:/csc606/image_assessment_augmented.csv', low_memory=False)

# convert int columns to str
labeled_df['leg']=labeled_df['leg'].astype(str)
labeled_df['site']=labeled_df['site'].astype(str)
labeled_df['hole']=labeled_df['hole'].astype(str)
labeled_df['core']=labeled_df['core'].astype(str)
labeled_df['section']=labeled_df['section'].astype(str)

# Create single column in dataframe called "key" that concatenates leg, site, hole,
# The key column will be used for joining/merging later
labeled_df['key']=labeled_df['leg'] + "." + labeled_df['site'] + "." + labeled_df['
```

In [7]:

```
labeled_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 131 entries, 0 to 130
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
---  -- 
 0   leg               131 non-null    object  
 1   site              131 non-null    object  
 2   hole              131 non-null    object  
 3   core              131 non-null    object  
 4   section           131 non-null    object  
 5   file_name         131 non-null    object  
 6   greater_than_50_percent_bad  131 non-null    bool    
 7   notes             58 non-null     object  
 8   key               131 non-null    object  
dtypes: bool(1), object(8)
memory usage: 8.4+ KB
```

```
In [8]: labeled_df
```

Out[8]:

| | leg | site | hole | core | section | | file_name | greater_th |
|------------|------------|-------------|-------------|-------------|----------------|--|--|-------------------|
| 0 | 166 | 1006 | A | 13 | 4 | | leg166_site1006_holeA_core13_section4.png | |
| 1 | 167 | 1010 | B | 1 | 3 | | leg167_site1010_holeB_core1_section3.png | |
| 2 | 167 | 1010 | C | 4 | 3 | | leg167_site1010_holeC_core4_section3.png | |
| 3 | 167 | 1010 | E | 10 | 4 | | leg167_site1010_holeE_core10_section4.png | |
| 4 | 167 | 1010 | E | 9 | 3 | | leg167_site1010_holeE_core9_section3.png | |
| ... | ... | ... | ... | ... | ... | | ... | |
| 126 | 398 | U1591 | B | 5 | 4 | | leg398_siteU1591_holeB_core5_section4.png | |
| 127 | 398 | U1593 | A | 34 | 1 | | leg398_siteU1593_holeA_core34_section1.png | |
| 128 | 398 | U1598 | A | 8 | 1 | | leg398_siteU1598_holeA_core8_section1.png | |
| 129 | 398 | U1600 | B | 19 | 1 | | leg398_siteU1600_holeB_core19_section1.png | |
| 130 | 398 | U1600 | B | 5 | 1 | | leg398_siteU1600_holeB_core5_section1.png | |

131 rows × 9 columns

In [9]: `# Here we merge the Labeled dataset with the curated dataset
merged_labeled_df = pd.merge(labeled_df, df, on='key', how='left')`

In [10]: `merged_labeled_df`

Out[10]:

| | leg_x | site_x | hole_x | core_x | section_x | file_name | gr |
|-------------|--------------|---------------|---------------|---------------|------------------|---|-----------|
| 0 | 166 | 1006 | A | 13 | 4 | leg166_site1006_holeA_core13_section4.png | |
| 1 | 166 | 1006 | A | 13 | 4 | leg166_site1006_holeA_core13_section4.png | |
| 2 | 166 | 1006 | A | 13 | 4 | leg166_site1006_holeA_core13_section4.png | |
| 3 | 166 | 1006 | A | 13 | 4 | leg166_site1006_holeA_core13_section4.png | |
| 4 | 166 | 1006 | A | 13 | 4 | leg166_site1006_holeA_core13_section4.png | |
| ... | ... | ... | ... | ... | ... | ... | |
| 7015 | 398 | U1600 | B | 5 | 1 | leg398_siteU1600_holeB_core5_section1.png | |
| 7016 | 398 | U1600 | B | 5 | 1 | leg398_siteU1600_holeB_core5_section1.png | |
| 7017 | 398 | U1600 | B | 5 | 1 | leg398_siteU1600_holeB_core5_section1.png | |
| 7018 | 398 | U1600 | B | 5 | 1 | leg398_siteU1600_holeB_core5_section1.png | |
| 7019 | 398 | U1600 | B | 5 | 1 | leg398_siteU1600_holeB_core5_section1.png | |

7020 rows × 29 columns

In [11]:

```
# here we create the features: mean/mode/median/std/min/max for both depth and com
# we also create the 25%(Q1) and 75%(Q3) quantiles for compressional velocity withi
#
# added in v2: mean/mode/median/std/min/max for the calculated rho feature
#
groupby_columns = ['key','greater_than_50_percent_bad']
df_grouped = merged_labeled_df.groupby(groupby_columns)[['depth_m','compressional_v
    depth_mean=('depth_m','mean'),
    depth_median=('depth_m','median'),
    depth_mode=('depth_m',lambda x: x.mode()[0]),
    depth_std=('depth_m', 'std'),
    depth_min=('depth_m', 'min'),
    depth_max=('depth_m', 'max'),
    velocity_mean=('compressional_velocity(m/s)', 'mean'),
    velocity_median=('compressional_velocity(m/s)', 'median'),
    velocity_mode=('compressional_velocity(m/s)', lambda x: x.mode()[0]),
    velocity_std=('compressional_velocity(m/s)', 'std'),
    velocity_min=('compressional_velocity(m/s)', 'min'),
    velocity_max=('compressional_velocity(m/s)', 'max'),
    velocity_q1 =('compressional_velocity(m/s)',lambda x: x.quantile(0.25)),
    velocity_q3 =('compressional_velocity(m/s)',lambda x: x.quantile(0.75)),
    rho_mean=('rho', 'mean'),
    rho_median=('rho', 'median'),
    rho_mode=('rho',lambda x: x.mode()[0]),
    rho_std=('rho', 'std'),
```

```
    rho_min=('rho','min'),  
    rho_max=('rho','max')  
)
```

In [12]: df_grouped

Out[12]:

| | | | depth_mean | depth_median | depth_mode |
|------------------|-----|-----------------------------|------------|--------------|------------|
| | key | greater_than_50_percent_bad | | | |
| 166.1006.A.13.4 | | False | 116.851433 | 116.8715 | 116.150 |
| 167.1010.B.1.3 | | False | 3.480000 | 3.4800 | 3.030 |
| 167.1010.C.4.3 | | False | 28.210000 | 28.2100 | 27.530 |
| 167.1010.E.10.4 | | True | 90.240000 | 90.2400 | 89.550 |
| 167.1010.E.9.3 | | False | 79.240000 | 79.2400 | 78.550 |
| | ... | ... | ... | ... | ... |
| 398.U1591.B.5.4 | | False | 40.932646 | 40.9330 | 40.354 |
| 398.U1593.A.34.1 | | True | 227.806071 | 227.8125 | 227.100 |
| 398.U1598.A.8.1 | | False | 61.359717 | 61.3500 | 60.700 |
| 398.U1600.B.19.1 | | False | 87.416433 | 87.4240 | 86.700 |
| 398.U1600.B.5.1 | | True | 21.648475 | 21.6480 | 20.900 |

131 rows × 20 columns



In [13]: # calculate the interquartile range for compressional velocity between Q1 and Q3
calculate upper and lower boundaries for compressional velocity which are 1.5x the
IGR and 1.5x the IGR below the median
df_grouped['velocity_igr'] = df_grouped['velocity_q3'] - df_grouped['velocity_q1']
df_grouped['velocity_upper'] = df_grouped['velocity_median'] + (df_grouped['velocit
df_grouped['velocity_lower'] = df_grouped['velocity_median'] - (df_grouped['velocit

In [14]: df_grouped

Out[14]:

| | | | depth_mean | depth_median | depth_mode |
|-----|-------------------------|-----------------------------|------------|--------------|------------|
| | key | greater_than_50_percent_bad | | | |
| 1 | 166.1006.A.13.4 | False | 116.851433 | 116.8715 | 116.150 |
| 2 | 167.1010.B.1.3 | False | 3.480000 | 3.4800 | 3.030 |
| 3 | 167.1010.C.4.3 | False | 28.210000 | 28.2100 | 27.530 |
| 4 | 167.1010.E.10.4 | True | 90.240000 | 90.2400 | 89.550 |
| 5 | 167.1010.E.9.3 | False | 79.240000 | 79.2400 | 78.550 |
| ... | ... | ... | ... | ... | ... |
| 131 | 398.U1591.B.5.4 | False | 40.932646 | 40.9330 | 40.354 |
| 132 | 398.U1593.A.34.1 | True | 227.806071 | 227.8125 | 227.100 |
| 133 | 398.U1598.A.8.1 | False | 61.359717 | 61.3500 | 60.700 |
| 134 | 398.U1600.B.19.1 | False | 87.416433 | 87.4240 | 86.700 |
| 135 | 398.U1600.B.5.1 | True | 21.648475 | 21.6480 | 20.900 |

131 rows × 23 columns



In [15]:

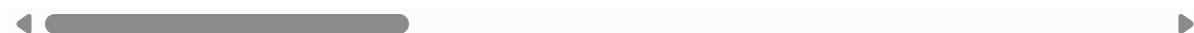
```
# get rid of the indexes created by groupby
df_grouped = df_grouped.reset_index()

# convert the boolean Label to numeric 1 and 0
df_grouped['label'] = df_grouped['greater_than_50_percent_bad'].astype(int)
df_grouped
```

Out[15]:

| | key | greater_than_50_percent_bad | depth_mean | depth_median | depth_mode |
|-----|------------------|-----------------------------|------------|--------------|------------|
| 0 | 166.1006.A.13.4 | False | 116.851433 | 116.8715 | 116.1! |
| 1 | 167.1010.B.1.3 | False | 3.480000 | 3.4800 | 3.0! |
| 2 | 167.1010.C.4.3 | False | 28.210000 | 28.2100 | 27.5! |
| 3 | 167.1010.E.10.4 | True | 90.240000 | 90.2400 | 89.5! |
| 4 | 167.1010.E.9.3 | False | 79.240000 | 79.2400 | 78.5! |
| ... | ... | ... | ... | ... | ... |
| 126 | 398.U1591.B.5.4 | False | 40.932646 | 40.9330 | 40.3! |
| 127 | 398.U1593.A.34.1 | True | 227.806071 | 227.8125 | 227.10 |
| 128 | 398.U1598.A.8.1 | False | 61.359717 | 61.3500 | 60.70 |
| 129 | 398.U1600.B.19.1 | False | 87.416433 | 87.4240 | 86.70 |
| 130 | 398.U1600.B.5.1 | True | 21.648475 | 21.6480 | 20.90 |

131 rows × 26 columns



```
In [16]: # pull out the key identifiers to X_identifiers, since they are text and can't be a
# drop the label columns and the key
X_identifiers = df_grouped['key']
X = df_grouped.drop('greater_than_50_percent_bad', axis=1).drop('key', axis=1).drop

# create the Labeled series from the label column
y = df_grouped['label']
```

```
In [17]: # Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [18]: # Model selection of Logistic Regression
model = LogisticRegression(max_iter=50000)
```

```
In [19]: # Model Training
model.fit(X_train, y_train)
```

Out[19]:

▾ LogisticRegression ⓘ ?
 LogisticRegression(max_iter=50000)

```
In [20]: # Model prediction using the test dataset
y_pred = model.predict(X_test)
```

```
In [21]: # Calculate prediction probabilities from test dataset
y_pred_prob = model.predict_proba(X_test)[:, 1]
```

```
In [22]: y_pred_prob
```

```
Out[22]: array([2.46717416e-04, 9.94993379e-01, 3.58020699e-01, 1.25775943e-02,
   3.42142605e-02, 4.81295229e-02, 1.13957797e-01, 7.96371636e-02,
   6.17198489e-02, 1.21676697e-01, 1.19821646e-01, 9.74089621e-02,
   5.86016347e-01, 4.24463806e-03, 1.62010579e-01, 1.03997401e-01,
   4.41887674e-02, 2.55587338e-02, 1.74499097e-01, 1.42091907e-01,
   1.92719303e-01, 8.82836554e-02, 9.49209263e-20, 4.17216757e-01,
   3.72961391e-01, 3.38447666e-02, 9.27158845e-02])
```

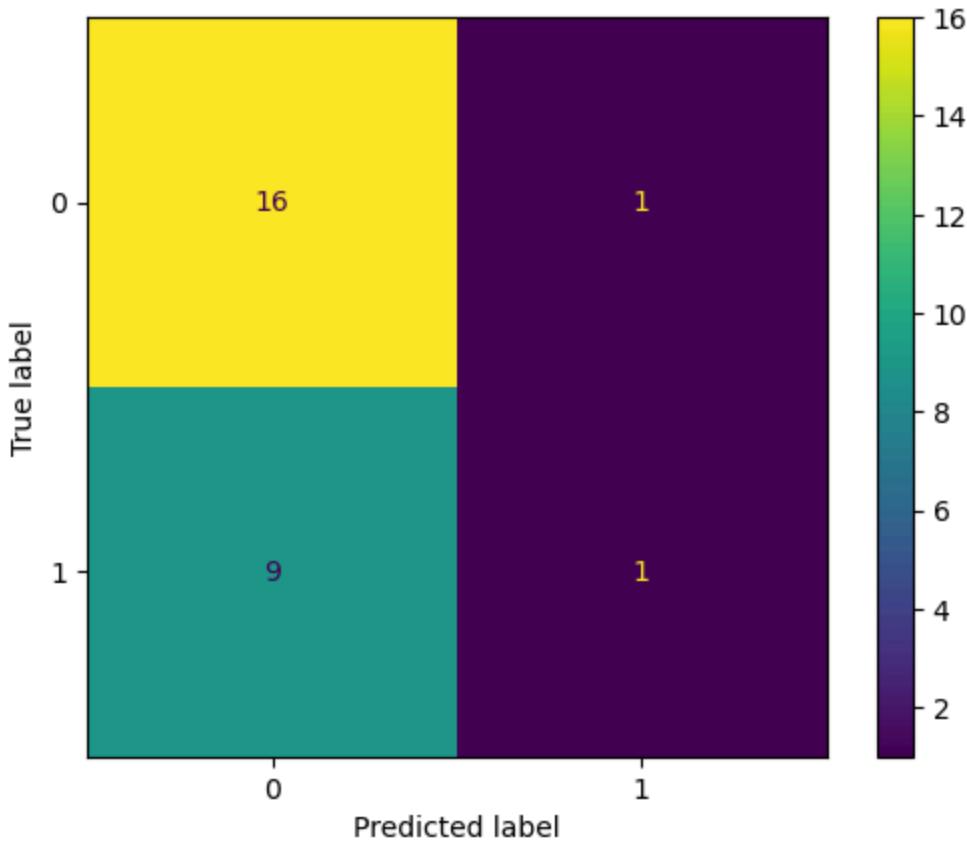
```
In [23]: # Calculate accuracy and create a classification report having precision, recall, f
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
```

```
In [24]: print(f"Accuracy: {accuracy:.2f}")
print("\nClassification Report:\n", report)
```

Accuracy: 0.63

| Classification Report: | | | | |
|------------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.64 | 0.94 | 0.76 | 17 |
| 1 | 0.50 | 0.10 | 0.17 | 10 |
| accuracy | | | 0.63 | 27 |
| macro avg | 0.57 | 0.52 | 0.46 | 27 |
| weighted avg | 0.59 | 0.63 | 0.54 | 27 |

```
In [25]: # build confusion matrix for the test dataset
conf_mat = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_mat,
                               display_labels=model.classes_)
disp.plot()
plt.show()
```



```
In [26]: fs = SelectKBest(score_func=mutual_info_regression, k=4) # k=4 based on results of
fs.fit(X_train, y_train)
selected_features_mask = fs.get_support()
print(X_train.columns[selected_features_mask])
```

```
Index(['velocity_mean', 'velocity_median', 'velocity_q3', 'velocity_upper'], dtype='object')
```

```
In [27]: x_train_fs = fs.transform(X_train)
x_test_fs = fs.transform(X_test)
```

```
In [28]: x_train_fs.shape
```

```
Out[28]: (104, 4)
```

```
In [29]: model2 = LogisticRegression(max_iter=50000)
model2.fit(x_train_fs, y_train)
model2.score(x_test_fs, y_test)
```

```
Out[29]: 0.7037037037037037
```

```
In [30]: model2.score(x_train_fs, y_train)
```

```
Out[30]: 0.9038461538461539
```

```
In [31]: # Plot the importance of each feature
# Code from: Murach's Python for Data Science, 2nd Edition, Chapter 11
df1 = pd.DataFrame(X_train.columns[selected_features_mask], columns=['feature'])
```

```

df2 = pd.DataFrame(fs.scores_[selected_features_mask], columns=['importance'])
importance = df1.join(df2)
importance.sort_values('importance', ascending=False).head()

```

Out[31]:

| | feature | importance |
|----------|-----------------|------------|
| 2 | velocity_q3 | 0.295706 |
| 1 | velocity_median | 0.239710 |
| 3 | velocity_upper | 0.239058 |
| 0 | velocity_mean | 0.232435 |

In [32]:

```

# Score the model for test/train datasets for varying number of features
#
# Code from: Murach's Python for Data Science, 2nd Edition, Chapter 11
#

model2 = LogisticRegression()
testScores = []
trainScores = []

for i in range(1, len(X_train.columns)):
    fs = SelectKBest(score_func=mutual_info_regression, k=i)
    fs.fit(X_train, y_train)

    x_train_fs = fs.transform(X_train)
    x_test_fs = fs.transform(X_test)

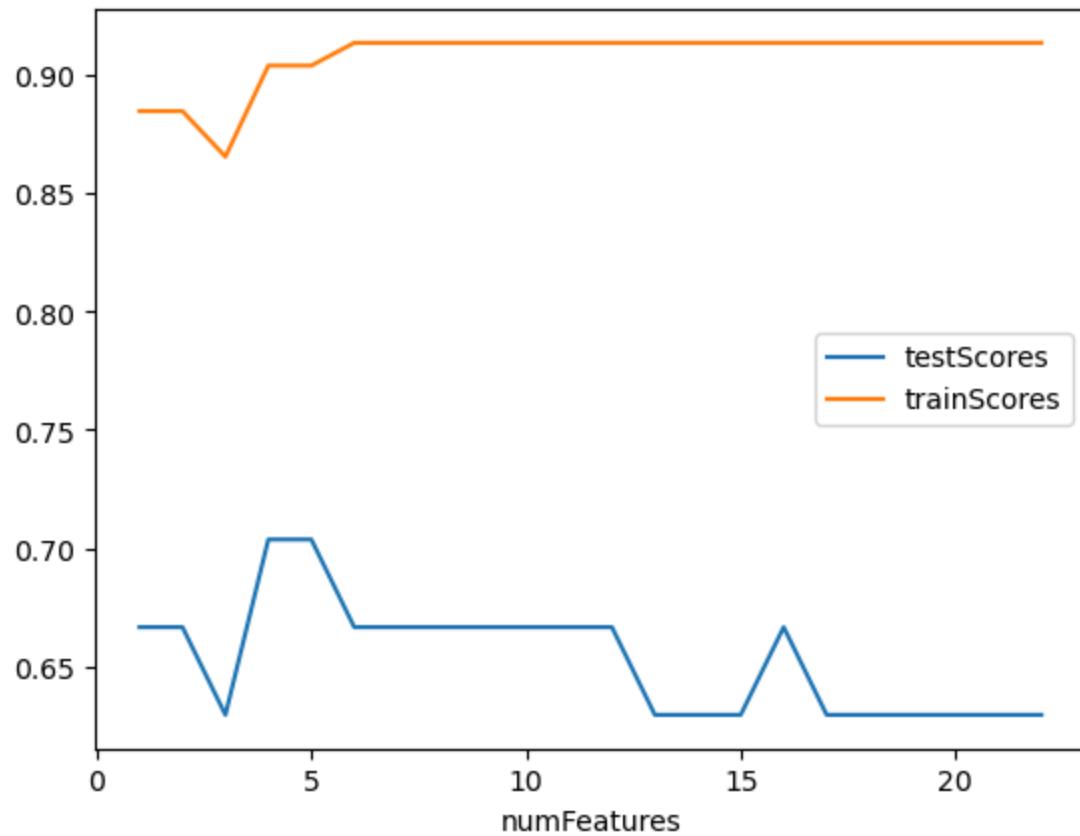
    model.fit(x_train_fs, y_train)

    testScore = model.score(x_test_fs, y_test)
    trainScore = model.score(x_train_fs, y_train)
    testScores.append(testScore)
    trainScores.append(trainScore)

df = pd.DataFrame(data={'testScores':testScores, 'trainScores':trainScores})
df.reset_index(inplace=True)
df.rename(columns={'index':'numFeatures'}, inplace=True)
df.numFeatures = df.numFeatures + 1
df.plot(x='numFeatures', y=['testScores', 'trainScores'])

```

Out[32]: <Axes: xlabel='numFeatures'>



In []:

In []: